

C:\cygwin\home\zahorjan\cse461\12au\05-retransmission.cp3

CSE 461 – Retransmission


Reliability

- Two strategies to handle errors:
 - Detect and retransmit, or Automatic Repeat reQuest. (ARQ)
 - Error correcting codes, or *Forward Error Correction* (FEC)
- We've seen error detecting/correcting codes, now look at retransmissions



Operational Assumptions

- Frames can simply be undetected by receiver
 - It hears nothing at all
 - Therefore the receiver doesn't do anything in this situation
- Bits in frames can be mangled
 - Receiver detects, corrects when it can, but there are still errors
 - What should receiver do?
- Frames that are received (haven't gone undetected) are received in the order they were sent
 - No overtaking

Basic Operation



Retransmissions, or more formally Automatic Repeat Request (ARQ)



- Sender automatically resends after a timeout until a positive acknowledgment (ACK) is obtained from the receiver
- Receiver automatically acknowledges frames (packets) that are not corrupted or lost in the network
- ARQ is generic name for protocols based on this strategy

Two issues

1. How long to set the timeout?

- Only easy on a direct link, otherwise timing variability
 - Way too long lowers performance
 - Short implies sometimes timeout will be early

2. How to avoid accepting duplicate frames as new


- Given retransmissions, frame loss, and imprecise timeouts

Why might the receiver see duplicates of a single frame?


Timeouts

Retransmission timeout depends on round-trip time

- To send frame and receive an acknowledgement
- In general, need to account for variance on complex paths




LAN case – small,
regular RTT



Internet case –
large, varied RTT


CSSE by Tanenbaum & Wetherall, © Pearson Education-Prenice Hall and D. Wetherall, 2011

Problem cases (due to loss, timeouts)




- In the case of ACK loss (or poor choice of timeout) the receiver can't distinguish current message from next

The Need for Sequence Numbers




- Frame sequence numbers let receiver tell next frame from duplicate transmission

ACKs need sequence numbers too



The Problem Scenario




The Solution

- Hm, these things can be tricky!

Stop-and-Wait

- Only one outstanding frame at a time, 0 or 1.
- Retransmissions re-sent with same number
- Number only needs to distinguish between current and next frame
 - A single bit will do




Limitation of Stop-and-Wait



- Lousy performance if transmission time \ll prop. delay
 - How bad? You do the math
- Want to utilize all available bandwidth
 - Need to keep more data “in flight”
 - How much? The “bandwidth-delay product”:
 $\text{bits/sec} * \text{seconds} = \text{bits}$
- Leads to *Sliding Window Protocol*


Solution: Allow Multiple Frames in Flight

- This is a form of pipelining




Sliding Window Protocol

- There is some maximum number of un-ACK'ed frames the sender is allowed to have in flight
 - We call this “the window size”
 - Example: window size = 2



Sliding Window: Sender


- Assign sequence number to each frame (**SeqNum**)
- Maintain three state variables:
 - send window size (**SWS**)
 - last acknowledgment received (**LAR**)
 - last frame sent (**LFS**)
- Maintain invariant: **LFS - LAR ≤ SWS**



- Advance **LAR** when ACK arrives
- Buffer up to **sWS** frames

Sliding Window: Receiver

- Maintain three state variables
 - receive window size (**RWS**)
 - largest frame acceptable (**LFA**)
 - last frame received (**LFR**)
- Maintain invariant: $\mathbf{LFA - LFR \leq RWS}$




- Frame **seqNum** arrives:
 - if $\mathbf{LFR < seqNum \leq LFA} \Rightarrow$ accept else discard
 - send ACK to tell sender what has arrived (new or repeat)
- Advance **LFR** (and pass to application) as in-order frames arrive
- Need to buffer up to **RWS** frames

Acknowledgement options


- Different options are possible:
- Send cumulative ACKs – send ACK for largest frame such that all frames less than this have been received
 - Robust to ACK loss but not packet loss
- Send individual ACKs
 - Robust to packet loss but not ACK loss!
- Can combine:
 - Idea is to tell the sender what frames the receiver already has
 - Usually have cumulative ACK plus hints

Sliding Window Example



Sequence Number Space

- SeqNum field is finite; sequence numbers wrap around
- Sequence number space must be larger than number of outstanding frames
- $SWS \leq \text{MaxSeqNum} - 1$ is not sufficient
- $SWS < (\text{MaxSeqNum} + 1) / 2$ is correct rule
- Intuitively, SeqNum “slides” between two halves of sequence number space



Sliding Window Summary

- It is perhaps the best known algorithm in networking
- First role is to enable reliable delivery of packets
 - Timeouts and acknowledgements
 - This has been our focus
- Second role is to enable in order delivery of packets
 - Receiver doesn't pass data up to app until it has packets in order
- Third role is to enable pipelined transmission
 - Crucial for high latency transmissions
- Fourth role is to enable flow control
 - Prevents fast sender from overflowing slow receiver's buffer
 - We will see this when we get to TCP

When to use ARQ or FEC?

- Will depend on the kind of errors and cost of recovery
- Example: Message with 1000 bits, Prob(bit error) 0.001
 - Case 1: random errors
 - Case 2: bursts of 1000 errors
- Q: What to use in Case 1 and 2?

ARQ vs. FEC

- FEC used at low-level to lower residual error rate
- ARQ often used to fix large errors, e.g., packet collision, and with detection to protect against residual errors
- FEC sometimes used at high level too:
 - Real time applications (no time to retransmit!)
 - Nice interaction with broadcast (different receiver errors!)

Example: 802.11

- The standard scheme is:
- PHY: FEC on data via interleaving and a binary convolutional code or LDPC
 - rates from $\frac{1}{2}$ to $\frac{5}{6}$.
- PHY header has 16 bit CRC
- Link: 32 bit CRC on frame and retransmission